

1. Library and data loading

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from scipy import stats
from scipy.stats import randint

# prep
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.datasets import make_classification
from sklearn.preprocessing import Binarizer, LabelEncoder,
MinMaxScaler

# models
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
ExtraTreesClassifier

# Validation libraries
from sklearn import metrics
from sklearn.metrics import accuracy_score, mean_squared_error,
precision_recall_curve
from sklearn.model_selection import cross_val_score

# Neural Network
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV

#Bagging
from sklearn.ensemble import BaggingClassifier, AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier

#Naive bayes
from sklearn.naive_bayes import GaussianNB

#Stacking
from mlxtend.classifier import StackingClassifier

#reading in CSV's from a file path
train_df = pd.read_csv('input\\survey.csv')
```

```
#Pandas: whats the data row count?
```

```
print("\n" , train_df.shape)
```

```
#Pandas: whats the distribution of the data?
```

```
print("\n" , train_df.describe())
```

```
#Pandas: What types of data do i have?
```

```
print("\n" , train_df.info())
```

```
(1259, 27)
```

```

              Age
count  1.259000e+03
mean   7.942815e+07
std    2.818299e+09
min    -1.726000e+03
25%    2.700000e+01
50%    3.100000e+01
75%    3.600000e+01
max     1.000000e+11
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1259 entries, 0 to 1258
```

```
Data columns (total 27 columns):
```

#	Column	Non-Null Count	Dtype
0	Timestamp	1259 non-null	object
1	Age	1259 non-null	int64
2	Gender	1259 non-null	object
3	Country	1259 non-null	object
4	state	744 non-null	object
5	self_employed	1241 non-null	object
6	family_history	1259 non-null	object
7	treatment	1259 non-null	object
8	work_interfere	995 non-null	object
9	no_employees	1259 non-null	object
10	remote_work	1259 non-null	object
11	tech_company	1259 non-null	object
12	benefits	1259 non-null	object
13	care_options	1259 non-null	object
14	wellness_program	1259 non-null	object
15	seek_help	1259 non-null	object
16	anonymity	1259 non-null	object
17	leave	1259 non-null	object
18	mental_health_consequence	1259 non-null	object
19	phys_health_consequence	1259 non-null	object
20	coworkers	1259 non-null	object
21	supervisor	1259 non-null	object
22	mental_health_interview	1259 non-null	object

23	phys_health_interview	1259 non-null	object
24	mental_vs_physical	1259 non-null	object
25	obs_consequence	1259 non-null	object
26	comments	164 non-null	object

dtypes: int64(1), object(26)

memory usage: 265.7+ KB

None

2. Data cleaning

#missing data

```
total = train_df.isnull().sum().sort_values(ascending=False)
```

```
percent =
```

```
(train_df.isnull().sum()/train_df.isnull().count()).sort_values(ascending=False)
```

```
missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
```

```
missing_data.head(20)
```

```
print(missing_data)
```

	Total	Percent
comments	1095	0.869738
state	515	0.409055
work_interfere	264	0.209690
self_employed	18	0.014297
seek_help	0	0.000000
obs_consequence	0	0.000000
mental_vs_physical	0	0.000000
phys_health_interview	0	0.000000
mental_health_interview	0	0.000000
supervisor	0	0.000000
coworkers	0	0.000000
phys_health_consequence	0	0.000000
mental_health_consequence	0	0.000000
leave	0	0.000000
anonymity	0	0.000000
Timestamp	0	0.000000
wellness_program	0	0.000000
Age	0	0.000000
benefits	0	0.000000
tech_company	0	0.000000
remote_work	0	0.000000
no_employees	0	0.000000
treatment	0	0.000000
family_history	0	0.000000
Country	0	0.000000

```
Gender          0  0.000000
care_options    0  0.000000
```

#dealing with missing data
#Removing the variables "Timestamp", "comments", "state" just to make our it easier to analyze.

```
train_df = train_df.drop(['comments'], axis= 1)
train_df = train_df.drop(['state'], axis= 1)
train_df = train_df.drop(['Timestamp'], axis= 1)
```

```
train_df.isnull().sum().max()
train_df.head(5)
```

	Age	Gender	Country	self_employed	family_history	treatment
0	37	Female	United States	NaN	No	Yes
1	44	M	United States	NaN	No	No
2	32	Male	Canada	NaN	No	No
3	31	Male	United Kingdom	NaN	Yes	Yes
4	31	Male	United States	NaN	No	No

	work_interfere	no_employees	remote_work	tech_company	...
0	Often	6-25	No	Yes	...
1	Rarely	More than 1000	No	No	... Don't know
2	Rarely	6-25	No	Yes	... Don't know
3	Often	26-100	No	Yes	...
4	Never	100-500	Yes	Yes	... Don't know

	leave mental_health_consequence
0	Somewhat easy No
1	Don't know Maybe
2	Somewhat difficult No
3	Somewhat difficult Yes
4	Don't know No

No

	coworkers	supervisor	mental_health_interview
phys_health_interview \			
0	Some of them	Yes	No
Maybe			
1	No	No	No
No			
2	Yes	Yes	Yes
Yes			
3	Some of them	No	Maybe
Maybe			
4	Some of them	Yes	Yes
Yes			

	mental_vs_physical	obs_consequence
0	Yes	No
1	Don't know	No
2	No	No
3	No	Yes
4	Don't know	No

[5 rows x 24 columns]

Cleaning NaN

Assign default values for each data type

```
defaultInt = 0
defaultString = 'NaN'
defaultFloat = 0.0
```

Create lists by data tpe

```
intFeatures = ['Age']
stringFeatures = ['Gender', 'Country', 'self_employed',
                  'family_history', 'treatment', 'work_interfere',
                  'no_employees', 'remote_work', 'tech_company',
                  'anonymity', 'leave', 'mental_health_consequence',
                  'phys_health_consequence', 'coworkers', 'supervisor',
                  'mental_health_interview', 'phys_health_interview',
                  'mental_vs_physical', 'obs_consequence', 'benefits',
                  'care_options', 'wellness_program',
                  'seek_help']
floatFeatures = []
```

Clean the NaN's

```
for feature in train_df:
    if feature in intFeatures:
        train_df[feature] = train_df[feature].fillna(defaultInt)
    elif feature in stringFeatures:
```

```

        train_df[feature] = train_df[feature].fillna(defaultString)
    elif feature in floatFeatures:
        train_df[feature] = train_df[feature].fillna(defaultFloat)
    else:
        print('Error: Feature %s not recognized.' % feature)
train_df.head(5)

```

	Age	Gender	Country	self_employed	family_history	treatment
0	37	Female	United States	NaN	No	Yes
1	44	M	United States	NaN	No	No
2	32	Male	Canada	NaN	No	No
3	31	Male	United Kingdom	NaN	Yes	Yes
4	31	Male	United States	NaN	No	No

	work_interfere	no_employees	remote_work	tech_company	...
0	Often	6-25	No	Yes	...
1	Rarely	More than 1000	No	No	... Don't know
2	Rarely	6-25	No	Yes	... Don't know
3	Often	26-100	No	Yes	...
4	Never	100-500	Yes	Yes	... Don't know

	leave mental_health_consequence
0	Somewhat easy No
1	Don't know Maybe
2	Somewhat difficult No
3	Somewhat difficult Yes
4	Don't know No

	coworkers supervisor mental_health_interview
0	Some of them Yes No

1	No	No	No
No			
2	Yes	Yes	Yes
Yes			
3	Some of them	No	Maybe
Maybe			
4	Some of them	Yes	Yes
Yes			

	mental_vs_physical	obs_consequence
0	Yes	No
1	Don't know	No
2	No	No
3	No	Yes
4	Don't know	No

[5 rows x 24 columns]

```
# clean 'Gender'
```

```
# lower case all column's elements
```

```
gender = train_df['Gender'].str.lower()
```

```
# Select unique elements
```

```
gender = train_df['Gender'].unique()
```

```
# Made gender groups
```

```
male_str = ["male", "m", "male-ish", "maile", "mal", "male (cis)",  
"make", "male ", "man", "msle", "mail", "malr", "cis man", "Cis Male",  
"cis male"]
```

```
trans_str = ["trans-female", "something kinda male?",  
"queer/she/they", "non-binary", "nah", "all", "enby", "fluid",  
"genderqueer", "androgynous", "agender", "male leaning androgynous",  
"guy (-ish) ^_^", "trans woman", "neuter", "female (trans)", "queer",  
"ostensibly male, unsure what that really means"]
```

```
female_str = ["cis female", "f", "female", "woman", "femake", "female",  
"cis-female/femme", "female (cis)", "femail"]
```

```
for (row, col) in train_df.iterrows():
```

```
    if str.lower(col.Gender) in male_str:  
        train_df['Gender'].replace(to_replace=col.Gender,  
value='male', inplace=True)
```

```
    if str.lower(col.Gender) in female_str:  
        train_df['Gender'].replace(to_replace=col.Gender,  
value='female', inplace=True)
```

```
    if str.lower(col.Gender) in trans_str:  
        train_df['Gender'].replace(to_replace=col.Gender,
```

```

value='trans', inplace=True)

stk_list = ['A little about you', 'p']
train_df = train_df[~train_df['Gender'].isin(stk_list)]

print(train_df['Gender'].unique())

['female' 'male' 'trans']

# complete missing age with mean
train_df['Age'].fillna(train_df['Age'].median(), inplace = True)

# Fill with media() values < 18 and > 120
s = pd.Series(train_df['Age'])
s[s<18] = train_df['Age'].median()
train_df['Age'] = s
s = pd.Series(train_df['Age'])
s[s>120] = train_df['Age'].median()
train_df['Age'] = s

# Ranges of Age
train_df['age_range'] = pd.cut(train_df['Age'], [0,20,30,65,100],
labels=["0-20", "21-30", "31-65", "66-100"], include_lowest=True)

#There are only 0.014% of self employed so let's change NaN to NOT
self_employed
#Replace "NaN" string from defaultString
train_df['self_employed'] =
train_df['self_employed'].replace([defaultString], 'No')
print(train_df['self_employed'].unique())

['No' 'Yes']

#There are only 0.20% of self work_interfere so let's change NaN to
"Don't know
#Replace "NaN" string from defaultString

train_df['work_interfere'] =
train_df['work_interfere'].replace([defaultString], 'Don\'t know' )
print(train_df['work_interfere'].unique())

['Often' 'Rarely' 'Never' 'Sometimes' "Don't know"]

```

3. Encoding data

```

#Encoding data
labelDict = {}
for feature in train_df:

```



```

le = preprocessing.LabelEncoder()
le.fit(train_df[feature])
le_name_mapping = dict(zip(le.classes_,
le.transform(le.classes_)))
train_df[feature] = le.transform(train_df[feature])
# Get labels
labelKey = 'label_' + feature
labelValue = [*le_name_mapping]
labelDict[labelKey] =labelValue

for key, value in labelDict.items():
    print(key, value)

```

#Get rid of 'Country'

```

train_df = train_df.drop(['Country'], axis= 1)
train_df.head()

```

```

label_Age [18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32,
33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49,
50, 51, 53, 54, 55, 56, 57, 58, 60, 61, 62, 65, 72]
label_Gender ['female', 'male', 'trans']
label_Country ['Australia', 'Austria', 'Belgium', 'Bosnia and
Herzegovina', 'Brazil', 'Bulgaria', 'Canada', 'China', 'Colombia',
'Costa Rica', 'Croatia', 'Czech Republic', 'Denmark', 'Finland',
'France', 'Georgia', 'Germany', 'Greece', 'Hungary', 'India',
'Ireland', 'Israel', 'Italy', 'Japan', 'Latvia', 'Mexico', 'Moldova',
'Netherlands', 'New Zealand', 'Nigeria', 'Norway', 'Philippines',
'Poland', 'Portugal', 'Romania', 'Russia', 'Singapore', 'Slovenia',
'South Africa', 'Spain', 'Sweden', 'Switzerland', 'Thailand', 'United
Kingdom', 'United States', 'Uruguay', 'Zimbabwe']
label_self_employed ['No', 'Yes']
label_family_history ['No', 'Yes']
label_treatment ['No', 'Yes']
label_work_interfere ["Don't know", 'Never', 'Often', 'Rarely',
'Sometimes']
label_no_employees ['1-5', '100-500', '26-100', '500-1000', '6-25',
'More than 1000']
label_remote_work ['No', 'Yes']
label_tech_company ['No', 'Yes']
label_benefits ["Don't know", 'No', 'Yes']
label_care_options ['No', 'Not sure', 'Yes']
label_wellness_program ["Don't know", 'No', 'Yes']
label_seek_help ["Don't know", 'No', 'Yes']
label_anonymity ["Don't know", 'No', 'Yes']
label_leave ["Don't know", 'Somewhat difficult', 'Somewhat easy',
'Very difficult', 'Very easy']
label_mental_health_consequence ['Maybe', 'No', 'Yes']
label_phys_health_consequence ['Maybe', 'No', 'Yes']
label_coworkers ['No', 'Some of them', 'Yes']

```

```

label_supervisor ['No', 'Some of them', 'Yes']
label_mental_health_interview ['Maybe', 'No', 'Yes']
label_phys_health_interview ['Maybe', 'No', 'Yes']
label_mental_vs_physical ["Don't know", 'No', 'Yes']
label_obs_consequence ['No', 'Yes']
label_age_range ['0-20', '21-30', '31-65', '66-100']

```

	Age	Gender	self_employed	family_history	treatment
work_interfere \					
0	19	0	0	0	1
2					
1	26	1	0	0	0
3					
2	14	1	0	0	0
3					
3	13	1	0	1	1
2					
4	13	1	0	0	0
1					

	no_employees	remote_work	tech_company	benefits	...	leave
0	4	0	1	2	...	2
1	5	0	0	0	...	0
2	4	0	1	1	...	1
3	2	0	1	1	...	1
4	1	1	1	2	...	0

	mental_health_consequence	phys_health_consequence	coworkers
supervisor \			
0	1	1	1
2			
1	0	1	0
0			
2	1	1	2
2			
3	2	2	1
0			
4	1	1	1
2			

	mental_health_interview	phys_health_interview	mental_vs_physical
\			
0	1	0	2
1	1	1	0
2	2	2	1
3	0	0	1

4		2	2	0
---	--	---	---	---

	obs_consequence	age_range
0	0	2
1	0	2
2	0	2
3	1	2
4	0	2

[5 rows x 24 columns]

Testing there aren't any missing data

```
#missing data
total = train_df.isnull().sum().sort_values(ascending=False)
percent = (train_df.isnull().sum()/train_df.isnull().count()).sort_values(ascending=False)
missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
missing_data.head(20)
print(missing_data)
```


	Total	Percent
Age	0	0.0
Gender	0	0.0
obs_consequence	0	0.0
mental_vs_physical	0	0.0
phys_health_interview	0	0.0
mental_health_interview	0	0.0
supervisor	0	0.0
coworkers	0	0.0
phys_health_consequence	0	0.0
mental_health_consequence	0	0.0
leave	0	0.0
anonymity	0	0.0
seek_help	0	0.0
wellness_program	0	0.0
care_options	0	0.0
benefits	0	0.0
tech_company	0	0.0
remote_work	0	0.0
no_employees	0	0.0
work_interfere	0	0.0
treatment	0	0.0
family_history	0	0.0
self_employed	0	0.0
age_range	0	0.0

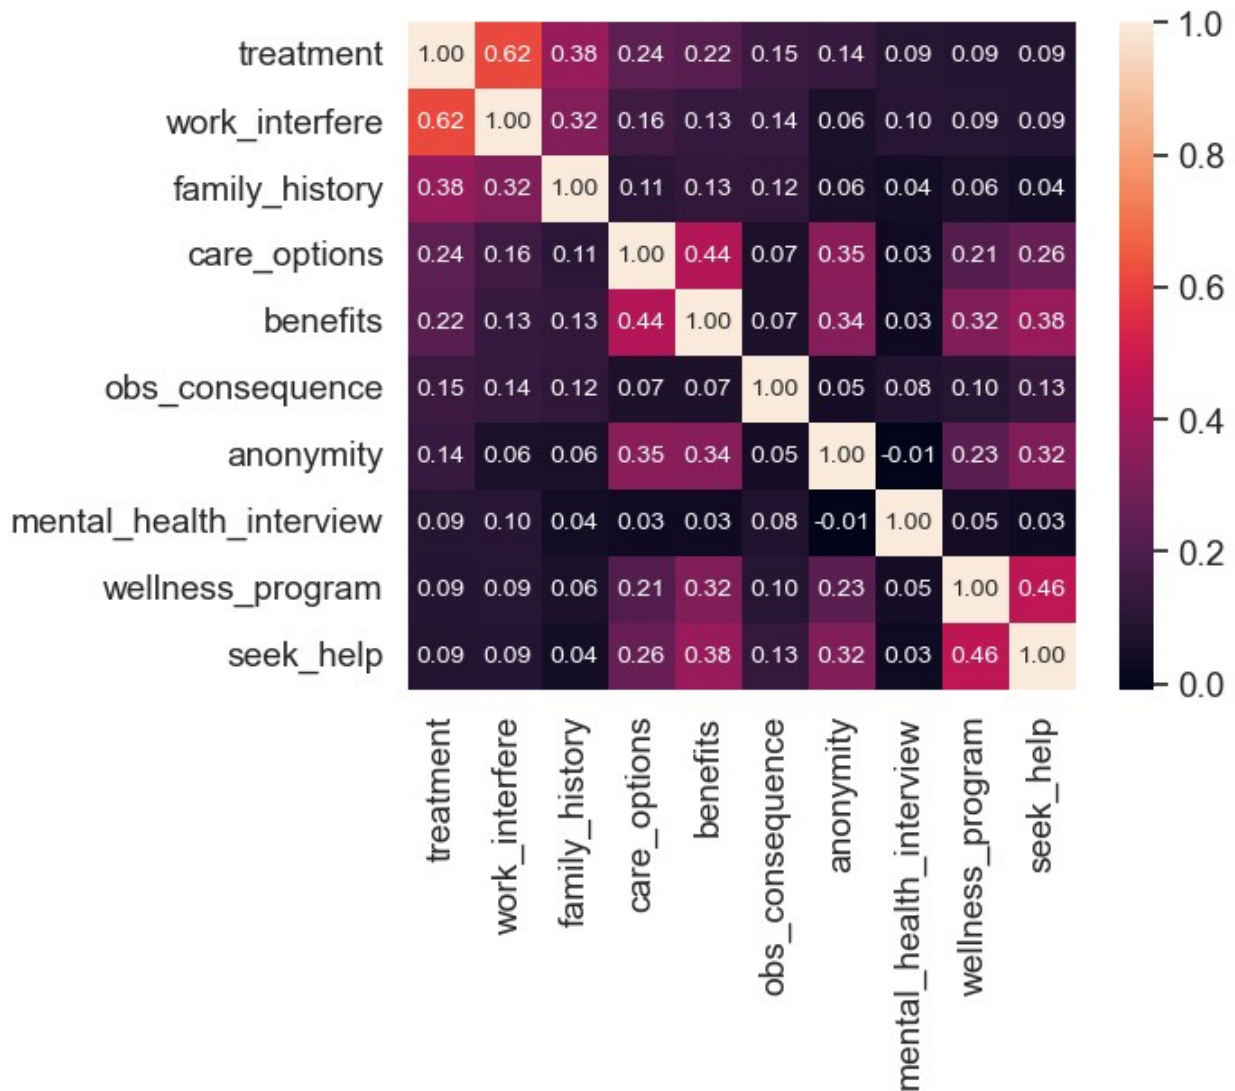
Features Scaling We're going to scale age, because is extremely different from the other ones.

4. Covariance Matrix. Variability comparison between categories of variables

```
#correlation matrix
corrmat = train_df.corr()
f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corrmat, vmax=.8, square=True);
plt.show()

#treatment correlation matrix
k = 10 #number of variables for heatmap
cols = corrmat.nlargest(k, 'treatment')['treatment'].index
cm = np.corrcoef(train_df[cols].values.T)
sns.set(font_scale=1.25)
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f',
annot_kws={'size': 10}, yticklabels=cols.values,
xticklabels=cols.values)
plt.show()
```





5. Some charts to see data relationship

Distribution and density by Age

```
# Distribution and density by Age
plt.figure(figsize=(12,8))
sns.distplot(train_df["Age"], bins=24)
plt.title("Distribution and density by Age")
plt.xlabel("Age")
```

C:\Users\athar\AppData\Local\Temp\ipykernel_10060\1394260443.py:3:
UserWarning:

`distplot` is a deprecated function and will be removed in seaborn


v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(train_df["Age"], bins=24)
```

```
Text(0.5, 0, 'Age')
```



Separate by treatment

```
# Separate by treatment or not
```

```
g = sns.FacetGrid(train_df, col='treatment', height=5)  
g = g.map(sns.distplot, "Age")
```

```
b:\Anaconda3\envs\_AIMLDataScience_env\Lib\site-packages\seaborn\  
axisgrid.py:848: UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
func(*plot_args, **plot_kwargs)
b:\Anaconda3\envs\_AIMLDataScience_env\Lib\site-packages\seaborn\
axisgrid.py:848: UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

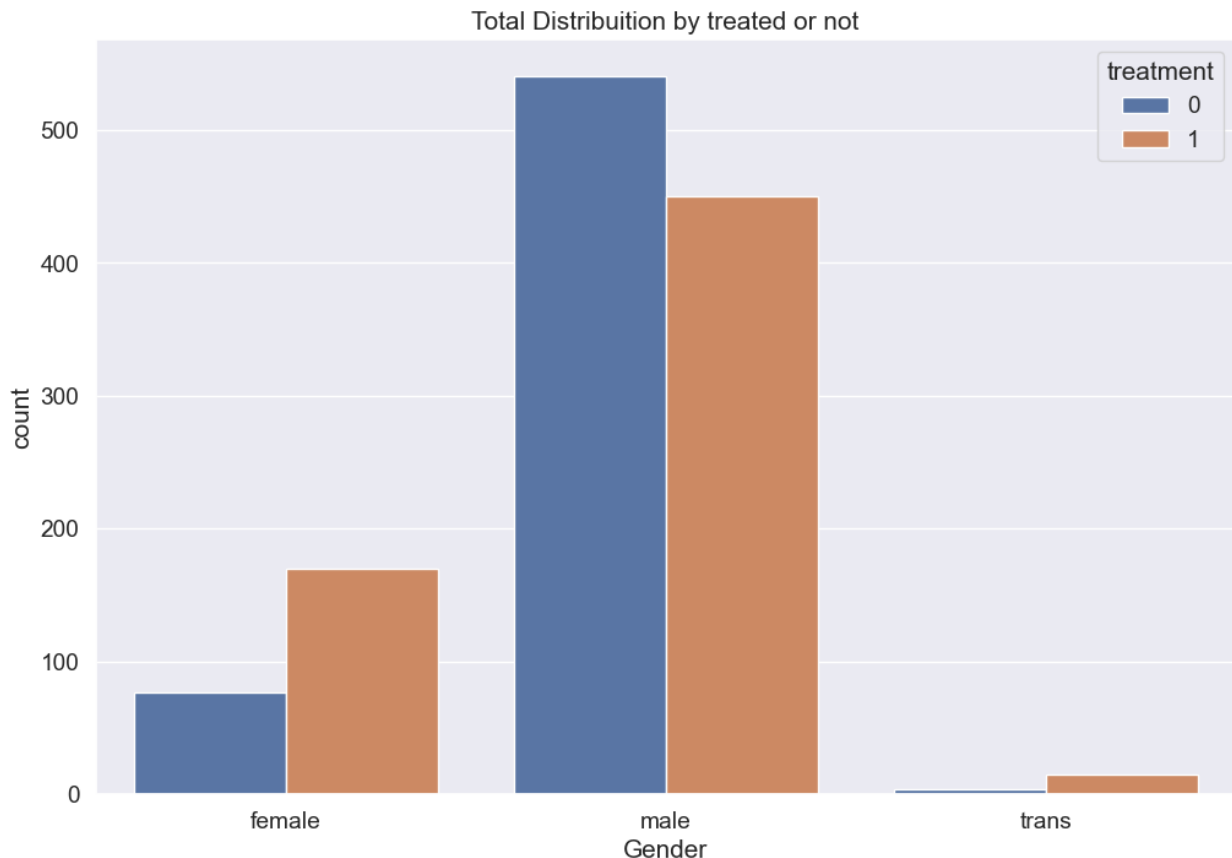
```
func(*plot_args, **plot_kwargs)
```



How many people has been treated?


```
# how many people has been treated
plt.figure(figsize=(12,8))
labels = labelDict['label_Gender']
g = sns.countplot(x="Gender", hue="treatment", data=train_df)
g.set_xticklabels(labels)

plt.title('Total Distribution by treated or not')
Text(0.5, 1.0, 'Total Distribution by treated or not')
```



Draw a nested barplot to show probabilities for class and sex

```
o = labelDict['label_age_range']

g = sns.catplot(x="age_range", y="treatment", hue="Gender",
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =
True)
g.set_xticklabels(o)

plt.title('Probability of mental health condition')
plt.ylabel('Probability x 100')
plt.xlabel('Age')
# replace legend labels
```

```
new_labels = labelDict['label_Gender']
for t, l in zip(g._legend.texts, new_labels): t.set_text(l)
```

```
# Positioning the legend
```


```
g.fig.subplots_adjust(top=0.9, right=0.8)
```

```
plt.show()
```

C:\Users\athar\AppData\Local\Temp\ipykernel_10060\827670349.py:3:
FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
g = sns.catplot(x="age_range", y="treatment", hue="Gender",
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =
True)
```



Barplot to show probabilities for family history

```
o = labelDict['label_family_history']
g = sns.catplot(x="family_history", y="treatment", hue="Gender",
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =
True)
```

```
g.set_xticklabels(o)
```

```
plt.title('Probability of mental health condition')
```

```
plt.ylabel('Probability x 100')
```

```
plt.xlabel('Family History')
```

```
# replace legend labels
```

```
new_labels = labelDict['label_Gender']
```

```
for t, l in zip(g._legend.texts, new_labels): t.set_text(l)
```


```
# Positioning the legend
g.fig.subplots_adjust(top=0.9, right=0.8)
```

```
plt.show()
```

C:\Users\athar\AppData\Local\Temp\ipykernel_10060\2317941586.py:2:
FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
g = sns.catplot(x="family_history", y="treatment", hue="Gender",
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =
True)
```



Barplot to show probabilities for care options

```
o = labelDict['label_care_options']
g = sns.catplot(x="care_options", y="treatment", hue="Gender",
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =
True)
g.set_xticklabels(o)
plt.title('Probability of mental health condition')
plt.ylabel('Probability x 100')
plt.xlabel('Care options')
```

```
# replace legend labels
new_labels = labelDict['label_Gender']
for t, l in zip(g._legend.texts, new_labels): t.set_text(l)
```


```
# Positioning the legend
```

```
g.fig.subplots_adjust(top=0.9, right=0.8)
plt.show()
```

C:\Users\athar\AppData\Local\Temp\ipykernel_10060\1780520276.py:2:
FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
g = sns.catplot(x="care_options", y="treatment", hue="Gender",
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =
True)
```



Barplot to show probabilities for benefits

```
o = labelDict['label_benefits']
g = sns.catplot(x="care_options", y="treatment", hue="Gender",
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =
True)
g.set_xticklabels(o)
plt.title('Probability of mental health condition')
plt.ylabel('Probability x 100')
plt.xlabel('Benefits')


# replace legend labels
new_labels = labelDict['label_Gender']
for t, l in zip(g._legend.texts, new_labels): t.set_text(l)

# Positioning the legend
g.fig.subplots_adjust(top=0.9, right=0.8)
plt.show()
```

C:\Users\athar\AppData\Local\Temp\ipykernel_10060\1084406037.py:2:
FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
g = sns.catplot(x="care_options", y="treatment", hue="Gender",  
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =  
True)
```



Barplot to show probabilities for work interfere


```
o = labelDict['label_work_interfere']  
g = sns.catplot(x="work_interfere", y="treatment", hue="Gender",  
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =  
True)  
g.set_xticklabels(o)  
plt.title('Probability of mental health condition')  
plt.ylabel('Probability x 100')  
plt.xlabel('Work interfere')  
  
# replace legend labels  
new_labels = labelDict['label_Gender']  
for t, l in zip(g._legend.texts, new_labels): t.set_text(l)  
  
# Positioning the legend  
g.fig.subplots_adjust(top=0.9, right=0.8)  
plt.show()
```

C:\Users\athar\AppData\Local\Temp\ipykernel_10060\2520480164.py:2:
FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same

effect.

```
g = sns.catplot(x="work_interfere", y="treatment", hue="Gender",
data=train_df, kind="bar", ci=None, height=5, aspect=2, legend_out =
True)
```



6. Scaling and fitting

Features Scaling We're going to scale age, because is extremely different from the others ones.

```
# Scaling Age
scaler = MinMaxScaler()
train_df['Age'] = scaler.fit_transform(train_df[['Age']])
train_df.head()
```

	Age	Gender	self_employed	family_history	treatment
work_interfere \					
0	0.431818	0	0	0	1
2					
1	0.590909	1	0	0	0
3					
2	0.318182	1	0	0	0
3					
3	0.295455	1	0	1	1
2					
4	0.295455	1	0	0	0
1					

	no_employees	remote_work	tech_company	benefits	...	leave	\
0	4	0	1	2	...	2	

1	5	0	0	0	...	0
2	4	0	1	1	...	1
3	2	0	1	1	...	1
4	1	1	1	2	...	0

	mental_health_consequence	phys_health_consequence	coworkers
supervisor \			
0	1	1	1
2			
1	0	1	0
0			
2	1	1	2
2			
3	2	2	1
0			
4	1	1	1
2			

	mental_health_interview	phys_health_interview	mental_vs_physical
\			
0	1	0	2
1	1	1	0
2	2	2	1
3	0	0	1
4	2	2	0

	obs_consequence	age_range
0	0	2
1	0	2
2	0	2
3	1	2
4	0	2

[5 rows x 24 columns]

Splitting the dataset

```
# define X and y
feature_cols = ['Age', 'Gender', 'family_history', 'benefits',
'care_options', 'anonymity', 'leave', 'work_interfere']
X = train_df[feature_cols]
y = train_df.treatment

# split X and y into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
```

```

test_size=0.30, random_state=0)

# Create dictionaries for final graph
# Use: methodDict['Stacking'] = accuracy_score
methodDict = {}
rmseDict = ()


# Build a forest and compute the feature importances
forest = ExtraTreesClassifier(n_estimators=250,
                              random_state=0)

forest.fit(X, y)
importances = forest.feature_importances_
std = np.std([tree.feature_importances_ for tree in
              forest.estimators_],
              axis=0)
indices = np.argsort(importances)[::-1]

labels = []
for f in range(X.shape[1]):
    labels.append(feature_cols[f])

# Plot the feature importances of the forest
plt.figure(figsize=(12,8))
plt.title("Feature importances")
plt.bar(range(X.shape[1]), importances[indices],
        color="r", yerr=std[indices], align="center")
plt.xticks(range(X.shape[1]), labels, rotation='vertical')
plt.xlim([-1, X.shape[1]])
plt.show()

```

7. Tuning

Evaluating a Classification Model

This function will evaluate:

- **Classification accuracy:** Percentage of correct predictions
- **Null accuracy:** Accuracy that could be achieved by always predicting the most frequent class
- **Percentage of ones**
- **Percentage of zeros**
- **Confusion matrix:** Table that describes the performance of a classification model
- **Confusion matrix elements:**

- True Positives (TP): Correctly predicted they have diabetes
- True Negatives (TN): Correctly predicted they don't have diabetes
- False Positives (FP): Incorrectly predicted they have diabetes (Type I error)
- False Negatives (FN): Incorrectly predicted they don't have diabetes (Type II error)
- **False Positive Rate**
- **Precision of Positive value**
- **AUC (Area Under the Curve):** Percentage of the ROC plot that is underneath the curve
 - .90-1 = excellent (A)
 - .80-.90 = good (B)
 - .70-.80 = fair (C)
 - .60-.70 = poor (D)
 - .50-.60 = fail (F)

And some other values for the tuning process.

```
def evalClassModel(model, y_test, y_pred_class, plot=False):
    #Classification accuracy: percentage of correct predictions
    # calculate accuracy
    print('Accuracy:', metrics.accuracy_score(y_test, y_pred_class))

    #Null accuracy: accuracy that could be achieved by always
    predicting the most frequent class
    # examine the class distribution of the testing set (using a
    Pandas Series method)
    print('Null accuracy:\n', y_test.value_counts())

    # calculate the percentage of ones
    print('Percentage of ones:', y_test.mean())

    # calculate the percentage of zeros
    print('Percentage of zeros:', 1 - y_test.mean())

    #Comparing the true and predicted response values
    print('True:', y_test.values[0:25])
    print('Pred:', y_pred_class[0:25])

    #Conclusion:
    #Classification accuracy is the easiest classification metric to
    understand
    #But, it does not tell you the underlying distribution of response
    values
    #And, it does not tell you what "types" of errors your classifier
    is making

    #Confusion matrix
```

```

# save confusion matrix and slice into four pieces
confusion = metrics.confusion_matrix(y_test, y_pred_class)
#[row, column]
TP = confusion[1, 1]
TN = confusion[0, 0]
FP = confusion[0, 1]
FN = confusion[1, 0]

# visualize Confusion Matrix
sns.heatmap(confusion,annot=True,fmt="d")
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

#Metrics computed from a confusion matrix
#Classification Accuracy: Overall, how often is the classifier
correct?
accuracy = metrics.accuracy_score(y_test, y_pred_class)
print('Classification Accuracy:', accuracy)

#Classification Error: Overall, how often is the classifier
incorrect?
print('Classification Error:', 1 - metrics.accuracy_score(y_test,
y_pred_class))

#False Positive Rate: When the actual value is negative, how often
is the prediction incorrect?
false_positive_rate = FP / float(TN + FP)
print('False Positive Rate:', false_positive_rate)

#Precision: When a positive value is predicted, how often is the
prediction correct?
print('Precision:', metrics.precision_score(y_test, y_pred_class))

# IMPORTANT: first argument is true values, second argument is
predicted probabilities
print('AUC Score:', metrics.roc_auc_score(y_test, y_pred_class))

# calculate cross-validated AUC
print('Cross-validated AUC:', cross_val_score(model, X, y, cv=10,
scoring='roc_auc').mean())

#####
#Adjusting the classification threshold
#####
# print the first 10 predicted responses
# 1D array (vector) of binary values (0, 1)

```

```

    print('First 10 predicted responses:\n', model.predict(X_test)
[0:10])

    # print the first 10 predicted probabilities of class membership
    print('First 10 predicted probabilities of class members:\n',
model.predict_proba(X_test)[0:10])

    # print the first 10 predicted probabilities for class 1
    model.predict_proba(X_test)[0:10, 1]

    # store the predicted probabilities for class 1
    y_pred_prob = model.predict_proba(X_test)[: , 1]

    if plot == True:
        # histogram of predicted probabilities
        # adjust the font size
        plt.rcParams['font.size'] = 12
        # 8 bins
        plt.hist(y_pred_prob, bins=8)

        # x-axis limit from 0 to 1
        plt.xlim(0,1)
        plt.title('Histogram of predicted probabilities')
        plt.xlabel('Predicted probability of treatment')
        plt.ylabel('Frequency')

    # predict treatment if the predicted probability is greater than
0.3
    # it will return 1 for all values above 0.3 and 0 otherwise
    # results are 2D so we slice out the first column
    y_pred_prob = y_pred_prob.reshape(-1,1)
    threshold = 0.3
    binarizer = Binarizer(threshold=threshold)
    y_pred_class = binarizer.transform(y_pred_prob)

    # print the first 10 predicted probabilities
    print('First 10 predicted probabilities:\n', y_pred_prob[0:10])

    #AUC is the percentage of the ROC plot that is underneath the
curve
    #Higher value = better classifier
    roc_auc = metrics.roc_auc_score(y_test, y_pred_prob)

    # IMPORTANT: first argument is true values, second argument is
predicted probabilities

```

```

# we pass y_test and y_pred_prob
# we do not use y_pred_class, because it will give incorrect
results without generating an error
# roc_curve returns 3 objects fpr, tpr, thresholds
# fpr: false positive rate
# tpr: true positive rate
fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_prob)
if plot == True:
    plt.figure()

    plt.plot(fpr, tpr, color='darkorange', label='ROC curve (area
= %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.rcParams['font.size'] = 12
    plt.title('ROC curve for treatment classifier')
    plt.xlabel('False Positive Rate (1 - Specificity)')
    plt.ylabel('True Positive Rate (Sensitivity)')
    plt.legend(loc="lower right")
    plt.show()

# define a function that accepts a threshold and prints
sensitivity and specificity
def evaluate_threshold(threshold):
    #Sensitivity: When the actual value is positive, how often is
the prediction correct?
    #Specificity: When the actual value is negative, how often is
the prediction correct?print('Sensitivity for ' + str(threshold) +
' :', tpr[thresholds > threshold][-1])
    print('Specificity for ' + str(threshold) + ' :', 1 -
fpr[thresholds > threshold][-1])

# One way of setting threshold
predict_mine = np.where(y_pred_prob > 0.50, 1, 0)
confusion = metrics.confusion_matrix(y_test, predict_mine)
print(confusion)

return accuracy

```

Tuning with cross validation score

```

#####
# Tuning with cross validation score
#####
def tuningCV(knn):

```

```

# search for an optimal value of K for KNN
k_range = list(range(1, 31))
k_scores = []
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X, y, cv=10, scoring='accuracy')
    k_scores.append(scores.mean())
print(k_scores)
# plot the value of K for KNN (x-axis) versus the cross-validated
accuracy (y-axis)
plt.plot(k_range, k_scores)
plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-Validated Accuracy')
plt.show()

```

Tuning with GridSearchCV

```

def tuningGridSerach(knn):
    #More efficient parameter tuning using GridSearchCV
    # define the parameter values that should be searched
    k_range = list(range(1, 31))
    print(k_range)

    # create a parameter grid: map the parameter names to the values
    that should be searched
    param_grid = dict(n_neighbors=k_range)
    print(param_grid)

    # instantiate the grid
    grid = GridSearchCV(knn, param_grid, cv=10, scoring='accuracy')

    # fit the grid with data
    grid.fit(X, y)

    # view the complete results (list of named tuples)
    grid.grid_scores_

    # examine the first tuple
    print(grid.grid_scores_[0].parameters)
    print(grid.grid_scores_[0].cv_validation_scores)
    print(grid.grid_scores_[0].mean_validation_score)

    # create a list of the mean scores only
    grid_mean_scores = [result.mean_validation_score for result in
    grid.grid_scores_]
    print(grid_mean_scores)

    # plot the results
    plt.plot(k_range, grid_mean_scores)

```

```

plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-Validated Accuracy')
plt.show()

# examine the best model
print('GridSearch best score', grid.best_score_)
print('GridSearch best params', grid.best_params_)
print('GridSearch best estimator', grid.best_estimator_)

```

Tuning with RandomizedSearchCV

```

def tuningRandomizedSearchCV(model, param_dist):
    #Searching multiple parameters simultaneously
    # n_iter controls the number of searches
    rand = RandomizedSearchCV(model, param_dist, cv=10,
scoring='accuracy', n_iter=10, random_state=5)
    rand.fit(X, y)
    rand.cv_results_

    # examine the best model
    print('Rand. Best Score: ', rand.best_score_)
    print('Rand. Best Params: ', rand.best_params_)

    # run RandomizedSearchCV 20 times (with n_iter=10) and record the
    best score
    best_scores = []
    for _ in range(20):
        rand = RandomizedSearchCV(model, param_dist, cv=10,
scoring='accuracy', n_iter=10)
        rand.fit(X, y)
        best_scores.append(round(rand.best_score_, 3))
    print(best_scores)

```

Tuning with searching multiple parameters simultaneously

```

def tuningMultParam(knn):
    #Searching multiple parameters simultaneously
    # define the parameter values that should be searched
    k_range = list(range(1, 31))
    weight_options = ['uniform', 'distance']

    # create a parameter grid: map the parameter names to the values
    that should be searched
    param_grid = dict(n_neighbors=k_range, weights=weight_options)
    print(param_grid)

    # instantiate and fit the grid
    grid = GridSearchCV(knn, param_grid, cv=10, scoring='accuracy')

```

```

grid.fit(X, y)

# view the complete results
print(grid.grid_scores_)

# examine the best model
print('Multiparam. Best Score: ', grid.best_score_)
print('Multiparam. Best Params: ', grid.best_params_)

```

8. Evaluating models

Logistic Regression

```

def logisticRegression():
    # train a logistic regression model on the training set
    logreg = LogisticRegression()
    logreg.fit(X_train, y_train)

    # make class predictions for the testing set
    y_pred_class = logreg.predict(X_test)

    print('##### Logistic Regression #####')


    accuracy_score = evalClassModel(logreg, y_test, y_pred_class,
    True)

    #Data for final graph
    methodDict['Log. Regres.'] = accuracy_score * 100

logisticRegression()


##### Logistic Regression #####
Accuracy: 0.7962962962962963
Null accuracy:
  treatment
0      191
1      187
Name: count, dtype: int64
Percentage of ones: 0.4947089947089947
Percentage of zeros: 0.5052910052910053
True: [0 0 0 0 0 0 0 0 1 1 0 1 1 0 1 1 0 0 0 1 1 0 0]
Pred: [1 0 0 0 1 1 0 1 0 1 0 1 1 0 1 1 1 0 0 0 0 1 0 0]


```

```
Classification Accuracy: 0.7962962962962963
Classification Error: 0.20370370370370372
False Positive Rate: 0.25654450261780104
Precision: 0.7644230769230769
AUC Score: 0.7968614385306716
Cross-validated AUC: 0.8753623882722146
First 10 predicted responses:
[1 0 0 0 1 1 0 1 0 1]
First 10 predicted probabilities of class members:
[[0.09193053 0.90806947]
 [0.95991564 0.04008436]
 [0.96547467 0.03452533]
 [0.78757121 0.21242879]
 [0.38959922 0.61040078]
 [0.05264207 0.94735793]
 [0.75035574 0.24964426]
 [0.19065116 0.80934884]
 [0.61612081 0.38387919]
 [0.47699963 0.52300037]]
First 10 predicted probabilities:
[[0.90806947]
 [0.04008436]]
```

```
[0.03452533]  
[0.21242879]  
[0.61040078]  
[0.94735793]  
[0.24964426]  
[0.80934884]  
[0.38387919]  
[0.52300037]]
```





```
[[142  49]
 [ 28 159]]
```

KNeighbors Classifier

```
def Knn():
    # Calculating the best parameters
    knn = KNeighborsClassifier(n_neighbors=5)

    # define the parameter values that should be searched
    k_range = list(range(1, 31))
    weight_options = ['uniform', 'distance']

    # specify "parameter distributions" rather than a "parameter grid"
    param_dist = dict(n_neighbors=k_range, weights=weight_options)
    tuningRandomizedSearchCV(knn, param_dist)

    # train a KNeighborsClassifier model on the training set
    knn = KNeighborsClassifier(n_neighbors=27, weights='uniform')
    knn.fit(X_train, y_train)
```

```

# make class predictions for the testing set
y_pred_class = knn.predict(X_test)

print('##### KNeighborsClassifier #####')

accuracy_score = evalClassModel(knn, y_test, y_pred_class, True)

#Data for final graph
methodDict['KNN'] = accuracy_score * 100

```


KNEIGHBORSCLASSIFIER

```

Knn()


Rand. Best Score: 0.8217650793650794
Rand. Best Params: {'weights': 'uniform', 'n_neighbors': 27}
[0.814, 0.819, 0.822, 0.811, 0.814, 0.814, 0.819, 0.819, 0.819, 0.822,
0.822, 0.819, 0.822, 0.817, 0.822, 0.819, 0.817, 0.816, 0.816, 0.813]
##### KNeighborsClassifier #####
Accuracy: 0.8042328042328042
Null accuracy:
  treatment
0      191
1      187
Name: count, dtype: int64
Percentage of ones: 0.4947089947089947
Percentage of zeros: 0.5052910052910053
True: [0 0 0 0 0 0 0 0 1 1 0 1 1 0 1 1 0 0 0 1 1 0 0]
Pred: [1 0 0 0 1 1 0 1 1 1 0 1 1 0 1 1 1 0 0 0 0 1 0 0]


```



```
Classification Accuracy: 0.8042328042328042
Classification Error: 0.1957671957671958
False Positive Rate: 0.2931937172774869
Precision: 0.7511111111111111
AUC Score: 0.8052747991152673
Cross-validated AUC: 0.8784644661702792
First 10 predicted responses:
[1 0 0 0 1 1 0 1 1 1]
First 10 predicted probabilities of class members:
[[0.33333333 0.66666667]
 [1.         0.         ]
 [1.         0.         ]
 [0.66666667 0.33333333]
 [0.37037037 0.62962963]
 [0.03703704 0.96296296]
 [0.59259259 0.40740741]
 [0.37037037 0.62962963]
 [0.33333333 0.66666667]
 [0.33333333 0.66666667]]
First 10 predicted probabilities:
[[0.66666667]
 [0.         ]
```

```
[0.      ]  
[0.33333333]  
[0.62962963]  
[0.96296296]  
[0.40740741]  
[0.62962963]  
[0.66666667]  
[0.66666667]
```





```
[[135  56]
 [ 18 169]]
```

Decision Tree classifier

```
def treeClassifier():
    # Calculating the best parameters
    tree = DecisionTreeClassifier()
    featuresSize = feature_cols.__len__()
    param_dist = {"max_depth": [3, None],
                  "max_features": randint(1, featuresSize),
                  "min_samples_split": randint(2, 9),
                  "min_samples_leaf": randint(1, 9),
                  "criterion": ["gini", "entropy"]}
    tuningRandomizedSearchCV(tree, param_dist)

    # train a decision tree model on the training set
    tree = DecisionTreeClassifier(max_depth=3, min_samples_split=8,
    max_features=6, criterion='entropy', min_samples_leaf=7)
    tree.fit(X_train, y_train)
```

```

# make class predictions for the testing set
y_pred_class = tree.predict(X_test)

print('##### Tree classifier #####')


accuracy_score = evalClassModel(tree, y_test, y_pred_class, True)

#Data for final graph
methodDict['Tree clas.'] = accuracy_score * 100

treeClassifier()


Rand. Best Score: 0.8305206349206349
Rand. Best Params: {'criterion': 'entropy', 'max_depth': 3,
'max_features': 6, 'min_samples_leaf': 7, 'min_samples_split': 8}
[0.831, 0.829, 0.831, 0.819, 0.831, 0.831, 0.821, 0.83, 0.831, 0.831,
0.831, 0.831, 0.831, 0.831, 0.816, 0.811, 0.831, 0.81, 0.831, 0.831]
##### Tree classifier #####
Accuracy: 0.8068783068783069
Null accuracy:
  treatment
0      191
1      187
Name: count, dtype: int64
Percentage of ones: 0.4947089947089947
Percentage of zeros: 0.5052910052910053
True: [0 0 0 0 0 0 0 0 1 1 0 1 1 0 1 1 0 0 0 1 1 0 0]
Pred: [1 0 0 0 1 1 0 1 1 1 0 1 1 0 1 1 1 0 0 0 0 1 0 0]


```

```
Classification Accuracy: 0.8068783068783069
Classification Error: 0.19312169312169314
False Positive Rate: 0.3193717277486911
Precision: 0.7415254237288136
AUC Score: 0.8082285746283282
Cross-validated AUC: 0.8845857468211298
First 10 predicted responses:
[1 0 0 0 1 1 0 1 1 1]
First 10 predicted probabilities of class members:
[[0.18823529 0.81176471]
 [0.95575221 0.04424779]
 [0.95575221 0.04424779]
 [0.95575221 0.04424779]
 [0.37583893 0.62416107]
 [0.05172414 0.94827586]
 [0.90384615 0.09615385]
 [0.37583893 0.62416107]
 [0.22018349 0.77981651]
 [0.22018349 0.77981651]]
First 10 predicted probabilities:
[[0.81176471]
 [0.04424779]
```

```
[0.04424779]  
[0.04424779]  
[0.62416107]  
[0.94827586]  
[0.09615385]  
[0.62416107]  
[0.77981651]  
[0.77981651]]
```





```
[[130  61]
 [ 12 175]]
```

Random Forests

```
def randomForest():
    # Calculating the best parameters
    forest = RandomForestClassifier(n_estimators = 20)

    featuresSize = feature_cols.__len__()
    param_dist = {"max_depth": [3, None],
                  "max_features": randint(1, featuresSize),
                  "min_samples_split": randint(2, 9),
                  "min_samples_leaf": randint(1, 9),
                  "criterion": ["gini", "entropy"]}
    tuningRandomizedSearchCV(forest, param_dist)

    # Building and fitting my_forest
    forest = RandomForestClassifier(max_depth = None,
    min_samples_leaf=8, min_samples_split=2, n_estimators = 20,
```

```

random_state = 1)
my_forest = forest.fit(X_train, y_train)

# make class predictions for the testing set
y_pred_class = my_forest.predict(X_test)

print('##### Random Forests #####')

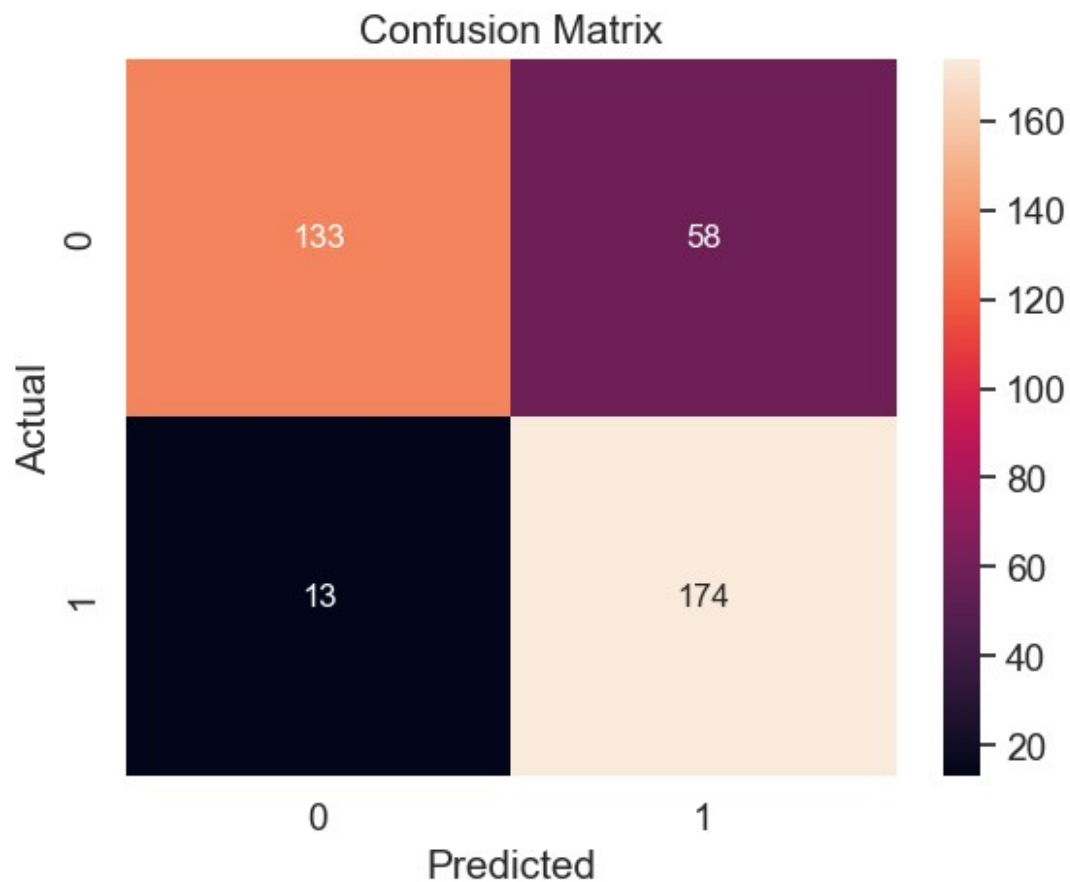
accuracy_score = evalClassModel(my_forest, y_test, y_pred_class,
True)

#Data for final graph
methodDict['R. Forest'] = accuracy_score * 100

randomForest()


Rand. Best Score: 0.8305206349206349
Rand. Best Params: {'criterion': 'entropy', 'max_depth': 3,
'max_features': 6, 'min_samples_leaf': 7, 'min_samples_split': 8}
[0.832, 0.831, 0.831, 0.833, 0.831, 0.833, 0.831, 0.831, 0.831, 0.831,
0.832, 0.831, 0.831, 0.831, 0.832, 0.831, 0.831, 0.831, 0.831, 0.831]
##### Random Forests #####
Accuracy: 0.8121693121693122
Null accuracy:
  treatment
0      191
1      187
Name: count, dtype: int64
Percentage of ones: 0.4947089947089947
Percentage of zeros: 0.5052910052910053
True: [0 0 0 0 0 0 0 0 1 1 0 1 1 0 1 1 0 0 0 1 1 0 0]
Pred: [1 0 0 0 1 1 0 1 1 1 0 1 1 0 1 1 1 0 0 0 0 1 0 0]


```



```
Classification Accuracy: 0.8121693121693122
Classification Error: 0.1878306878306878
False Positive Rate: 0.3036649214659686
Precision: 0.75
AUC Score: 0.8134081809782457
Cross-validated AUC: 0.8934280651104528
First 10 predicted responses:
[1 0 0 0 1 1 0 1 1 1]
First 10 predicted probabilities of class members:
[[0.2555794  0.7444206 ]
 [0.95069083 0.04930917]
 [0.93851009 0.06148991]
 [0.87096597 0.12903403]
 [0.40653554 0.59346446]
 [0.17282958 0.82717042]
 [0.89450448 0.10549552]
 [0.4065912  0.5934088 ]
 [0.20540631 0.79459369]
 [0.19337644 0.80662356]]
First 10 predicted probabilities:
[[0.7444206 ]
 [0.04930917]]
```

```
[0.06148991]  
[0.12903403]  
[0.59346446]  
[0.82717042]  
[0.10549552]  
[0.5934088 ]  
[0.79459369]  
[0.80662356]]
```





```
[[133  58]
 [ 13 174]]
```

Bagging

```
def bagging():
    # Building and fitting
    bag = BaggingClassifier(DecisionTreeClassifier(), max_samples=1.0,
max_features=1.0, bootstrap_features=False)
    bag.fit(X_train, y_train)

    # make class predictions for the testing set
    y_pred_class = bag.predict(X_test)

    print('##### Bagging #####')

    accuracy_score = evalClassModel(bag, y_test, y_pred_class, True)

    #Data for final graph
    methodDict['Bagging'] = accuracy_score * 100
```

```
bagging()
```

```
##### Bagging #####
```

```
Accuracy: 0.7962962962962963
```

```
Null accuracy:
```

```
  treatment
```

```
0      191
```

```
1      187
```


```
Name: count, dtype: int64
```

```
Percentage of ones: 0.4947089947089947
```

```
Percentage of zeros: 0.5052910052910053
```

```
True: [0 0 0 0 0 0 0 0 1 1 0 1 1 0 1 1 0 1 0 0 0 1 1 0 0]
```

```
Pred: [1 0 0 0 0 1 0 0 1 1 0 1 1 1 1 0 1 1 0 0 0 0 1 0 0]
```



```
Classification Accuracy: 0.7962962962962963
```

```
Classification Error: 0.20370370370370372
```

```
False Positive Rate: 0.24607329842931938
```

```
Precision: 0.7696078431372549
```

```
AUC Score: 0.7967494470420248
```

```
Cross-validated AUC: 0.8516097220698315
```

```
First 10 predicted responses:
```


```
[1 0 0 0 0 1 0 0 1 1]
```



First 10 predicted probabilities of class members:

```
[[0.33333333 0.66666667]
 [1.         0.         ]
 [1.         0.         ]
 [0.8        0.2        ]
 [0.8        0.2        ]
 [0.1        0.9        ]
 [1.         0.         ]
 [0.9        0.1        ]
 [0.         1.         ]
 [0.1        0.9        ]]
```

First 10 predicted probabilities:

```
[[0.66666667]
 [0.         ]
 [0.         ]
 [0.2        ]
 [0.2        ]
 [0.9        ]
 [0.         ]
 [0.1        ]
 [1.         ]
 [0.9        ]]
```





```
[[144  47]
 [ 30 157]]
```

Boosting

```
def boosting():
    # Building and fitting
    clf = DecisionTreeClassifier(criterion='entropy', max_depth=1)
    boost = AdaBoostClassifier(base_estimator=clf, n_estimators=500)
    boost.fit(X_train, y_train)

    # make class predictions for the testing set
    y_pred_class = boost.predict(X_test)

    print('##### Boosting #####')

    accuracy_score = evalClassModel(boost, y_test, y_pred_class, True)

    #Data for final graph
    methodDict['Boosting'] = accuracy_score * 100
```


```

boosting()

b:\Anaconda3\envs\_AIMLDataScience_env\Lib\site-packages\sklearn\
ensemble\_base.py:166: FutureWarning: `base_estimator` was renamed to
`estimator` in version 1.2 and will be removed in 1.4.
  warnings.warn(

##### Boosting #####
Accuracy: 0.8174603174603174
Null accuracy:
  treatment
0      191
1      187
Name: count, dtype: int64
Percentage of ones: 0.4947089947089947
Percentage of zeros: 0.5052910052910053
True: [0 0 0 0 0 0 0 0 1 1 0 1 1 0 1 1 0 1 0 0 0 1 1 0 0]
Pred: [1 0 0 0 0 1 0 1 1 1 0 1 1 0 1 1 1 1 0 0 0 0 1 0 0]

```



```

Classification Accuracy: 0.8174603174603174
Classification Error: 0.18253968253968256
False Positive Rate: 0.28272251308900526

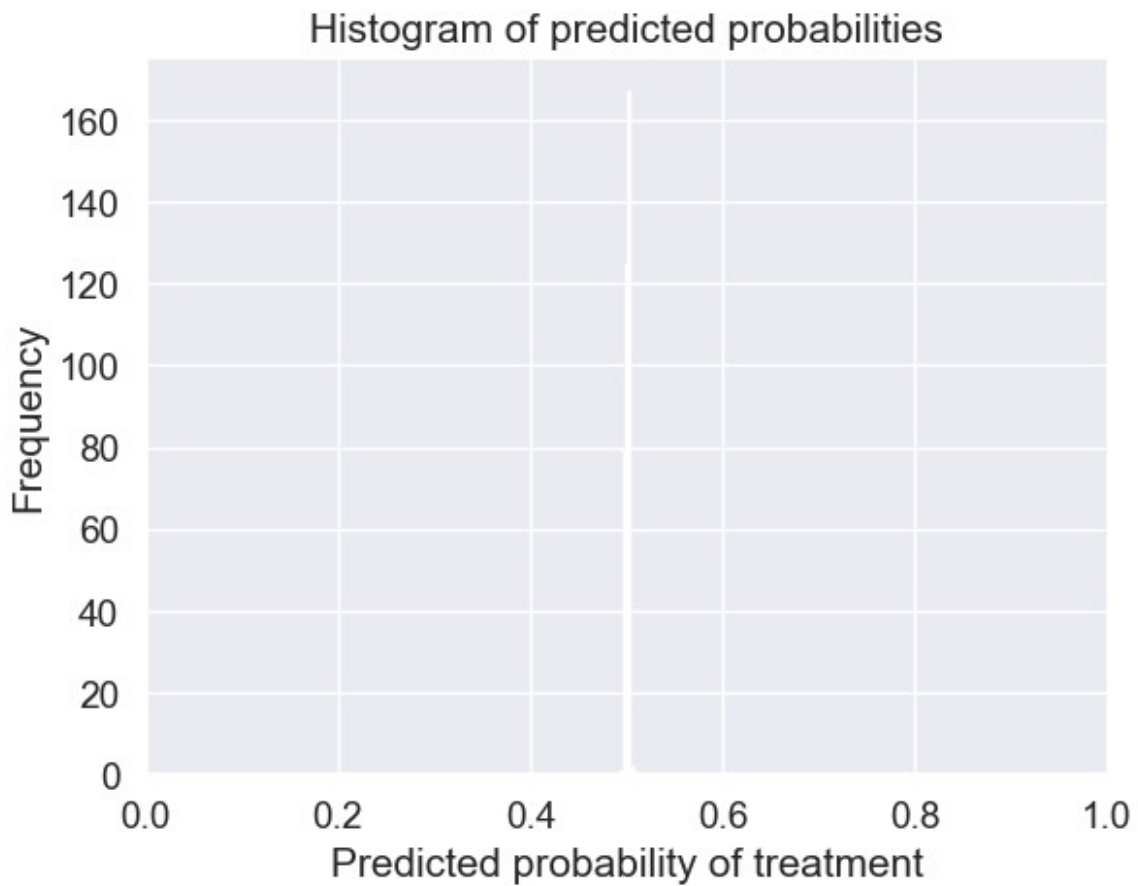
```


```
Precision: 0.7610619469026548
AUC Score: 0.8185317915838397
```

[illegible]

```
Cross-validated AUC: 0.8746279095195426
First 10 predicted responses:
[1 0 0 0 0 1 0 1 1 1]
First 10 predicted probabilities of class members:
[[0.49924555 0.50075445]
 [0.50285507 0.49714493]]
```

```
[0.50291786 0.49708214]
[0.50127788 0.49872212]
[0.50013552 0.49986448]
[0.49796157 0.50203843]
[0.50046371 0.49953629]
[0.49939483 0.50060517]
[0.49921757 0.50078243]
[0.49897133 0.50102867]]
First 10 predicted probabilities:
[[0.50075445]
 [0.49714493]
 [0.49708214]
 [0.49872212]
 [0.49986448]
 [0.50203843]
 [0.49953629]
 [0.50060517]
 [0.50078243]
 [0.50102867]]
```





```
[[137  54]
 [ 15 172]]
```

Stacking

```
def stacking():
    # Building and fitting
    clf1 = KNeighborsClassifier(n_neighbors=1)
    clf2 = RandomForestClassifier(random_state=1)
    clf3 = GaussianNB()
    lr = LogisticRegression()
    stack = StackingClassifier(classifiers=[clf1, clf2, clf3],
    meta_classifier=lr)
    stack.fit(X_train, y_train)

    # make class predictions for the testing set
    y_pred_class = stack.predict(X_test)

    print('##### Stacking #####')
```

```

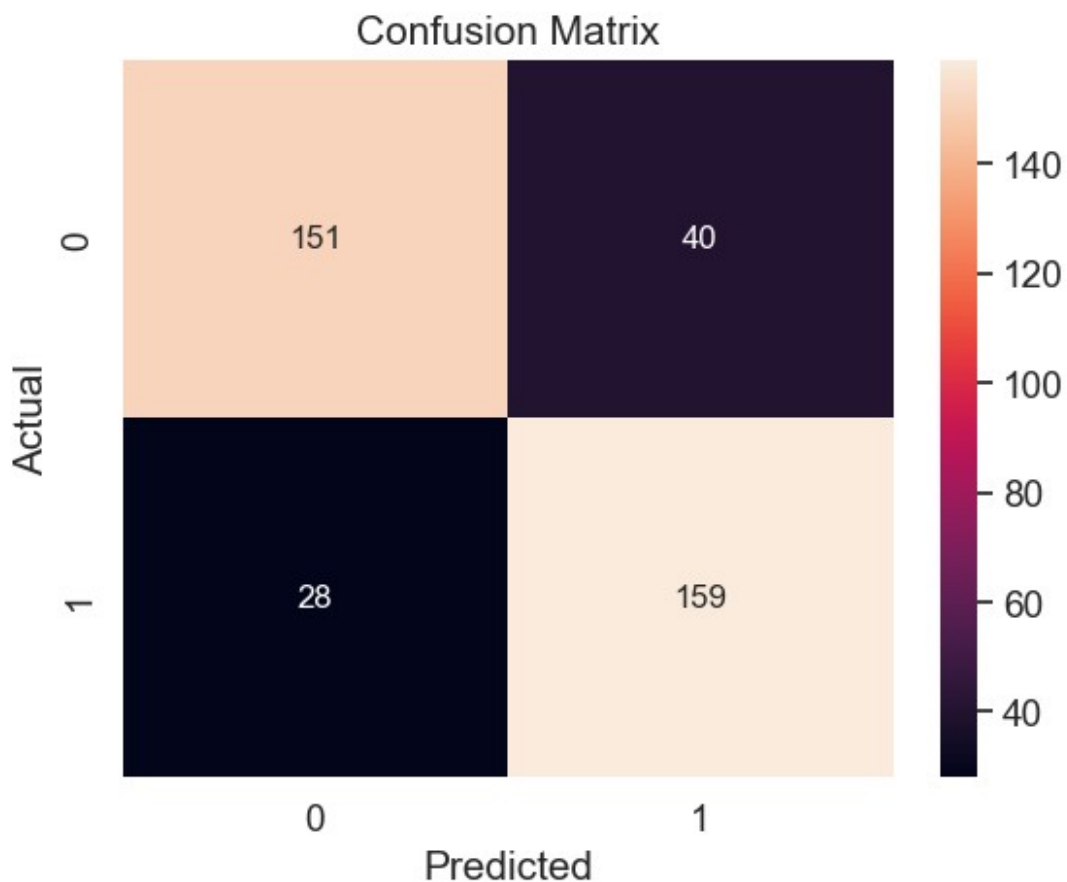
accuracy_score = evalClassModel(stack, y_test, y_pred_class, True)

#Data for final graph
methodDict['Stacking'] = accuracy_score * 100

stacking()

##### Stacking #####
Accuracy: 0.8201058201058201
Null accuracy:
  treatment
0      191
1      187
Name: count, dtype: int64
Percentage of ones: 0.4947089947089947
Percentage of zeros: 0.5052910052910053
True: [0 0 0 0 0 0 0 0 1 1 0 1 1 0 1 1 0 1 0 0 0 1 1 0 0]
Pred: [1 0 0 0 0 1 0 0 1 1 0 1 1 1 1 1 0 1 0 0 0 0 0 1 0 0]

```





```

Classification Accuracy: 0.8201058201058201
Classification Error: 0.17989417989417988

```

False Positive Rate: 0.2094240837696335
Precision: 0.7989949748743719
AUC Score: 0.8204216479547554
Cross-validated AUC: 0.8431811731188892
First 10 predicted responses:
[1 0 0 0 0 1 0 0 1 1]
First 10 predicted probabilities of class members:
[[0.01710346 0.98289654]
[0.98675465 0.01324535]
[0.98675465 0.01324535]
[0.98675465 0.01324535]
[0.98675465 0.01324535]
[0.01710346 0.98289654]
[0.98675465 0.01324535]
[0.97307936 0.02692064]
[0.03462234 0.96537766]
[0.01710346 0.98289654]]
First 10 predicted probabilities:
[[0.98289654]
[0.01324535]
[0.01324535]
[0.01324535]
[0.01324535]
[0.98289654]
[0.01324535]
[0.02692064]
[0.96537766]
[0.98289654]]





```
[[151  40]
 [ 28 159]]
```

9. Predicting with Neural Network

Create input functions

```
import tensorflow as tf
import argparse

batch_size = 100
train_steps = 1000

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.30, random_state=0)

def train_input_fn(features, labels, batch_size):
    """An input function for training"""
    # Convert the inputs to a Dataset.
```

```

    dataset = tf.data.Dataset.from_tensor_slices((dict(features),
labels))

    # Shuffle, repeat, and batch the examples.
    return dataset.shuffle(1000).repeat().batch(batch_size)

def eval_input_fn(features, labels, batch_size):
    """An input function for evaluation or prediction"""
    features=dict(features)
    if labels is None:
        # No labels, use only features.
        inputs = features
    else:
        inputs = (features, labels)

    # Convert the inputs to a Dataset.
    dataset = tf.data.Dataset.from_tensor_slices(inputs)

    # Batch the examples
    assert batch_size is not None, "batch_size must not be None"
    dataset = dataset.batch(batch_size)

    # Return the dataset.
    return dataset

```

Define the feature columns

A feature column is an object describing how the model should use raw input data from the features dictionary.

```

# Define Tensorflow feature columns
age = tf.feature_column.numeric_column("Age")
gender = tf.feature_column.numeric_column("Gender")
family_history = tf.feature_column.numeric_column("family_history")
benefits = tf.feature_column.numeric_column("benefits")
care_options = tf.feature_column.numeric_column("care_options")
anonymity = tf.feature_column.numeric_column("anonymity")
leave = tf.feature_column.numeric_column("leave")
work_interfere = tf.feature_column.numeric_column("work_interfere")
feature_columns = [age, gender, family_history, benefits,
care_options, anonymity, leave, work_interfere]

```

WARNING:tensorflow:From C:\Users\athar\AppData\Local\Temp\ipykernel_10060\3225071575.py:2: numeric_column (from tensorflow.python.feature_column.feature_column_v2) is deprecated and will be removed in a future version.
Instructions for updating:
Use Keras preprocessing layers instead, either directly or via the `tf.keras.utils.FeatureSpace` utility. Each of `tf.feature_column.*`

has a functional equivalent in `tf.keras.layers` for feature preprocessing when training a Keras model.

Instantiate an Estimator

Our problem is a classic classification problem. We want to predict whether a patient has to be treated or not. We'll use `tf.estimator.DNNClassifier` for deep models that perform multi-class classification.

```
# Build a DNN with 2 hidden layers and 10 nodes in each hidden layer.
model = tf.estimator.DNNClassifier(feature_columns=feature_columns,
                                   hidden_units=[10, 10],

optimizer=tf.keras.optimizers.legacy.Adam(
    learning_rate=0.1
))
```

```
WARNING:tensorflow:From C:\Users\athar\AppData\Local\Temp\
ipykernel_10060\3956841515.py:2: DNNClassifierV2.__init__ (from
tensorflow_estimator.python.estimator.canned.dnn) is deprecated and
will be removed in a future version.
```

```
Instructions for updating:
```

```
Use tf.keras instead.
```

```
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\head\
head_utils.py:54: BinaryClassHead.__init__ (from
tensorflow_estimator.python.estimator.head.binary_class_head) is
deprecated and will be removed in a future version.
```

```
Instructions for updating:
```

```
Use tf.keras instead.
```

```
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\canned\dnn.py:759:
Estimator.__init__ (from
tensorflow_estimator.python.estimator.estimator) is deprecated and
will be removed in a future version.
```

```
Instructions for updating:
```

```
Use tf.keras instead.
```

```
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\estimator.py:1844:
RunConfig.__init__ (from
tensorflow_estimator.python.estimator.run_config) is deprecated and
will be removed in a future version.
```

```
Instructions for updating:
```

```
Use tf.keras instead.
```

```
INFO:tensorflow:Using default config.
```

```
WARNING:tensorflow:Using temporary folder as model directory: C:\
Users\athar\AppData\Local\Temp\tmpk7sjx813
```

```
INFO:tensorflow:Using config: {'_model_dir': 'C:\\Users\\athar\\
AppData\\Local\\Temp\\tmpk7sjx813', '_tf_random_seed': None,
```

```

'_save_summary_steps': 100, '_save_checkpoints_steps': None,
'_save_checkpoints_secs': 600, '_session_config':
allow_soft_placement: true
graph_options {
  rewrite_options {
    meta_optimizer_iterations: ONE
  }
}
, '_keep_checkpoint_max': 5, '_keep_checkpoint_every_n_hours': 10000,
'_log_step_count_steps': 100, '_train_distribute': None, '_device_fn':
None, '_protocol': None, '_eval_distribute': None,
'_experimental_distribute': None,
'_experimental_max_worker_delay_secs': None,
'_session_creation_timeout_secs': 7200, '_checkpoint_save_graph_def':
True, '_service': None, '_cluster_spec': ClusterSpec({}),
'_task_type': 'worker', '_task_id': 0, '_global_id_in_cluster': 0,
'_master': '', '_evaluation_master': '', '_is_chief': True,
'_num_ps_replicas': 0, '_num_worker_replicas': 1}

```

Train, Evaluate, and Predict

Now that we have an Estimator object, we can call methods to do the following:

- Train the model.
- Evaluate the trained model.
- Use the trained model to make predictions.

Train the model

The steps argument tells the method to stop training after a number of training steps.

```

model.train(input_fn=lambda:train_input_fn(X_train, y_train,
batch_size), steps=train_steps)

```

```

WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\estimator.py:385:
StopAtStepHook.__init__ (from
tensorflow.python.training.basic_session_run_hooks) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
INFO:tensorflow:Calling model_fn.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\model_fn.py:250:
EstimatorSpec.__new__ (from
tensorflow_estimator.python.estimator.model_fn) is deprecated and will
be removed in a future version.
Instructions for updating:
Use tf.keras instead.

```

```
INFO:tensorflow:Done calling model_fn.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\estimator.py:1416:
NanTensorHook.__init__ (from
tensorflow.python.training.basic_session_run_hooks) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\estimator.py:1419:
LoggingTensorHook.__init__ (from
tensorflow.python.training.basic_session_run_hooks) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow\python\training\
basic_session_run_hooks.py:232: SecondOrStepTimer.__init__ (from
tensorflow.python.training.basic_session_run_hooks) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\estimator.py:1456:
CheckpointSaverHook.__init__ (from
tensorflow.python.training.basic_session_run_hooks) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
INFO:tensorflow:Create CheckpointSaverHook.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow\python\training\monitored_session.py:579:
StepCounterHook.__init__ (from
tensorflow.python.training.basic_session_run_hooks) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow\python\training\monitored_session.py:586:
SummarySaverHook.__init__ (from
tensorflow.python.training.basic_session_run_hooks) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Calling checkpoint listeners before saving checkpoint
0...
```

```
INFO:tensorflow:Saving checkpoints for 0 into C:\Users\athar\AppData\Local\Temp\tmpk7sjx813\model.ckpt.
INFO:tensorflow:Calling checkpoint listeners after saving checkpoint 0...
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\site-packages\tensorflow\python\training\monitored_session.py:1455: SessionRunArgs.__new__ (from tensorflow.python.training.session_run_hook) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\site-packages\tensorflow\python\training\monitored_session.py:1454: SessionRunContext.__init__ (from tensorflow.python.training.session_run_hook) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\site-packages\tensorflow\python\training\monitored_session.py:1474: SessionRunValues.__new__ (from tensorflow.python.training.session_run_hook) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
INFO:tensorflow:loss = 0.6146792, step = 0
INFO:tensorflow:global_step/sec: 802.572
INFO:tensorflow:loss = 0.36849788, step = 100 (0.127 sec)
INFO:tensorflow:global_step/sec: 1349.14
INFO:tensorflow:loss = 0.31399375, step = 200 (0.073 sec)
INFO:tensorflow:global_step/sec: 1471.31
INFO:tensorflow:loss = 0.32438728, step = 300 (0.068 sec)
INFO:tensorflow:global_step/sec: 1467.63
INFO:tensorflow:loss = 0.43570128, step = 400 (0.068 sec)
INFO:tensorflow:global_step/sec: 1499.49
INFO:tensorflow:loss = 0.3671186, step = 500 (0.067 sec)
INFO:tensorflow:global_step/sec: 1485.53
INFO:tensorflow:loss = 0.38827735, step = 600 (0.068 sec)
INFO:tensorflow:global_step/sec: 1513.57
INFO:tensorflow:loss = 0.30714363, step = 700 (0.065 sec)
INFO:tensorflow:global_step/sec: 1418.35
INFO:tensorflow:loss = 0.31148607, step = 800 (0.071 sec)
INFO:tensorflow:global_step/sec: 1876.76
INFO:tensorflow:loss = 0.41515914, step = 900 (0.053 sec)
INFO:tensorflow:Calling checkpoint listeners before saving checkpoint 1000...
INFO:tensorflow:Saving checkpoints for 1000 into C:\Users\athar\AppData\Local\Temp\tmpk7sjx813\model.ckpt.
INFO:tensorflow:Calling checkpoint listeners after saving checkpoint
```

```
1000...
INFO:tensorflow:Loss for final step: 0.455875.

<tensorflow_estimator.python.estimator.canned.dnn.DNNClassifierV2 at
0x1ad02a78890>
```

Evaluate the trained model

Now that the model has been trained, we can get some statistics on its performance. The following code block evaluates the accuracy of the trained model on the test data.

```
# Evaluate the model.
eval_result = model.evaluate(
    input_fn=lambda:eval_input_fn(X_test, y_test, batch_size))

print('\nTest set accuracy: {accuracy:0.2f}\n'.format(**eval_result))

#Data for final graph
accuracy = eval_result['accuracy'] * 100
methodDict['NN DNNClasif.'] = accuracy

INFO:tensorflow:Calling model_fn.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Starting evaluation at 2023-11-07T09:15:55
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow\python\training\evaluation.py:260:
FinalOpsHook.__init__ (from
tensorflow.python.training.basic_session_run_hooks) is deprecated and
will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Restoring parameters from C:\Users\athar\AppData\
Local\Temp\tmpk7sjx813\model.ckpt-1000
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.
INFO:tensorflow:Inference Time : 0.38039s
INFO:tensorflow:Finished evaluation at 2023-11-07-09:15:55
INFO:tensorflow:Saving dict for global step 1000: accuracy =
0.7962963, accuracy_baseline = 0.505291, auc = 0.87921715,
auc_precision_recall = 0.849898, average_loss = 0.49389374,
global_step = 1000, label/mean = 0.49470899, loss = 0.49318892,
precision = 0.72540987, prediction/mean = 0.5608679, recall =
0.9465241
INFO:tensorflow:Saving 'checkpoint_path' summary for global step 1000:
C:\Users\athar\AppData\Local\Temp\tmpk7sjx813\model.ckpt-1000

Test set accuracy: 0.80
```


Making predictions (inferring) from the trained model

We now have a trained model that produces good evaluation results. We can now use the trained model to predict whether a patient needs treatment or not.

```
predictions =
list(model.predict(input_fn=lambda:eval_input_fn(X_train, y_train,
batch_size=batch_size)))

INFO:tensorflow:Calling model_fn.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\head\
base_head.py:786: ClassificationOutput.__init__ (from
tensorflow.python.saved_model.model_utils.export_output) is deprecated
and will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\head\
binary_class_head.py:561: RegressionOutput.__init__ (from
tensorflow.python.saved_model.model_utils.export_output) is deprecated
and will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
WARNING:tensorflow:From b:\Anaconda3\envs\_AIMLDataScience_env\Lib\
site-packages\tensorflow_estimator\python\estimator\head\
binary_class_head.py:563: PredictOutput.__init__ (from
tensorflow.python.saved_model.model_utils.export_output) is deprecated
and will be removed in a future version.
Instructions for updating:
Use tf.keras instead.
INFO:tensorflow:Done calling model_fn.
INFO:tensorflow:Graph was finalized.
INFO:tensorflow:Restoring parameters from C:\Users\athar\AppData\
Local\Temp\tmpk7sjx813\model.ckpt-1000
INFO:tensorflow:Running local_init_op.
INFO:tensorflow:Done running local_init_op.

# Generate predictions from the model
template = ('\nIndex: "{}", Prediction is "{}" ( {:.1f}%), expected
"{}"')

# Dictionary for predictions
col1 = []
col2 = []
col3 = []

for idx, input, p in zip(X_train.index, y_train, predictions):
    v = p["class_ids"][0]
```

```

class_id = p['class_ids'][0]
probability = p['probabilities'][class_id] # Probability

# Adding to dataframe
col1.append(idx) # Index
col2.append(v) # Prediction
col3.append(input) # Expecter

#print(template.format(idx, v, 100 * probability, input))

results = pd.DataFrame({'index':col1, 'prediction':col2,
'expected':col3})
results.head()

```

	index	prediction	expected
0	929	0	0
1	901	1	1
2	579	1	1
3	367	1	1
4	615	1	1


10. Success method plot

```

def plotSuccess():
    s = pd.Series(methodDict)
    s = s.sort_values(ascending=False)
    plt.figure(figsize=(12,8))
    #Colors
    ax = s.plot(kind='bar')
    for p in ax.patches:
        ax.annotate(str(round(p.get_height(),2)), (p.get_x() * 1.005,
p.get_height() * 1.005))
    plt.ylim([70.0, 90.0])
    plt.xlabel('Method')
    plt.ylabel('Percentage')
    plt.title('Success of methods')

    plt.show()
plotSuccess()

```



11. Creating predictions on test set

```
# Generate predictions with the best method
clf = AdaBoostClassifier()
clf.fit(X, y)
dfTestPredictions = clf.predict(X_test)

# Write predictions to csv file
# We don't have any significant field so we save the index
results = pd.DataFrame({'Index': X_test.index, 'Treatment':
dfTestPredictions})
# Save to file
# This file will be visible after publishing in the output section
results.to_csv('results.csv', index=False)
print(results)
```

	Index	Treatment
0	5	1

1	494	0
2	52	0
3	984	0
4	186	0
..
373	1084	1
374	506	0
375	1142	0
376	1124	0
377	689	1

[378 rows x 2 columns]